

Benjamin Philipose

206-816-9055 | philiposebenjamin@gmail.com | linkedin.com/in/bphilipose/ | github.com/bphilipose1

EDUCATION

Seattle University	Seattle, WA
<i>M.S. Computer Science, Specialization in Data Science (GPA: 4.00/4.00)</i>	Sep. 2023 – Jun. 2025
• Coursework: Large Language Models, Parallel Computing, Distributed Systems, Big Data, Machine Learning, AI, LLMs.	
Seattle University	Seattle, WA
<i>B.S. Computer Engineering & Computer Science (GPA: 4.00/4.00)</i>	
• Computer Science: Algorithms, Computing Systems, Databases, Object-Oriented Development, Web Development.	
• Computer Engineering: Embedded Systems, Machine Learning, Data Networks, Signals, Circuits, Microprocessor Design.	

EXPERIENCE

Machine Learning Researcher	06/2024 – 08/2024
<i>Seattle University</i>	Seattle, WA
• Developed deep learning models using PyTorch and CUDA for fMRI-based autism detection improving classification accuracy by 4%.	
• Designed a data pipeline for processing structured and unstructured time-series brain signal data.	
• Explored Graph Neural Networks (GNNs) and Transformer-based architectures for feature extraction and classification.	
Robotics Software Engineer	09/2023 – 06/2024
<i>Amazon Web Services (AWS)</i>	Seattle, WA
• Developed a semi-autonomous robotic cart integrating LiDAR and thermal imaging for real-time data center monitoring.	
• Optimized sensor data processing pipelines in Python and C++, improving data collection efficiency.	
• Implemented a barcode-scanning system using OpenCV for automated inventory management.	
• Led a 9-person engineering team, driving Agile development and AWS collaboration.	
Software Engineering Researcher	06/2023 – 08/2023
<i>National Institute of Standards and Technology (NIST)</i>	Gaithersburg, MD
• Designed and developed a high-performance, low-latency data transfer tool in C++ for autonomous vehicle simulations, enabling seamless multi-simulation connectivity.	
• Integrated NS-3, ROS2, and proprietary simulators to optimize simulation performance.	
• Collaborated with cross-functional teams to debug sensing platforms and refine real-time system performance.	
• Presented research on simulation optimization at the 2023 NIST Colloquium. Link to Article .	

PROJECTS

CUDA-Optimized Sorting & Scanning (GPU Acceleration, CUDA)	03/2024 – 04/2024
• Developed a CUDA-optimized parallel sorting and prefix scan algorithm using bitonic sort and dissemination prefix scan.	
• Implemented multi-tier GPU optimizations and profiled performance using NVIDIA Nsight Systems.	
• Achieved significant speedup over CPU-based methods by optimizing memory access and kernel launches.	
INRIX x AWS Hackathon Winner - ResponSight	12/2023
• Led a team to design a scalable ML model predicting traffic collision hotspots, improving EMS response times through enhanced data insights.	
• Developed a full-stack application using AWS SageMaker and Python for predictive analytics, deployed for EMS use cases. Link to Article .	
LegalEase: AI-Powered Business Chatbot (LLM, RAG, AI Deployment)	01/2025 – Present
• Developed a chatbot leveraging fine-tuned LLMs and Retrieval-Augmented Generation (RAG) to provide business filing recommendations.	
• Integrated legal document processing and ensured accuracy by grounding responses in statutory data.	
• Deployed on a secure cloud infrastructure to protect user data and optimize inference speed.	
Distributed Clustering with MPI (C++, HPC)	03/2024 – 04/2024
• Implemented a distributed K-Means clustering algorithm using MPI to process large datasets efficiently.	
• Optimized parallel execution by distributing centroid updates across 32 MPI processes, reducing computation time.	
• Applied clustering to high-dimensional datasets, achieving scalable and efficient pattern recognition.	
Parallelized Prefix Sum (C++, Multi-threading)	03/2024 – 04/2024
• Developed a multi-threaded prefix sum algorithm using Pthreads, achieving nearly 2x speedup over a sequential approach.	
• Optimized thread synchronization to reduce overhead and improve memory efficiency.	

SKILLS

Programming: Python, C++, CUDA, MATLAB, Bash, JavaScript
Machine Learning: PyTorch, TensorFlow, Scikit-Learn, OpenCV, GNNs, NLP, Transformers, RAG
HPC & Distributed Systems: CUDA, MPI, Parallel Computing, Nsight, Multithreading, Spark, Hadoop
Data Science: Pandas, NumPy, SciPy, Matplotlib, Seaborn, SQL
Cloud & Deployment: AWS (SageMaker, Lambda, EC2, S3), GCP, Docker, Kubernetes
Software Tools: Git, Linux, Jupyter Notebook, REST APIs, Unit Testing, Agile
AWS Cloud Practitioner (Link)